

Module 4 - Trees

SIT320 – Advanced Algorithms

Dr. Nayyar Zaidi

Motivations for BST

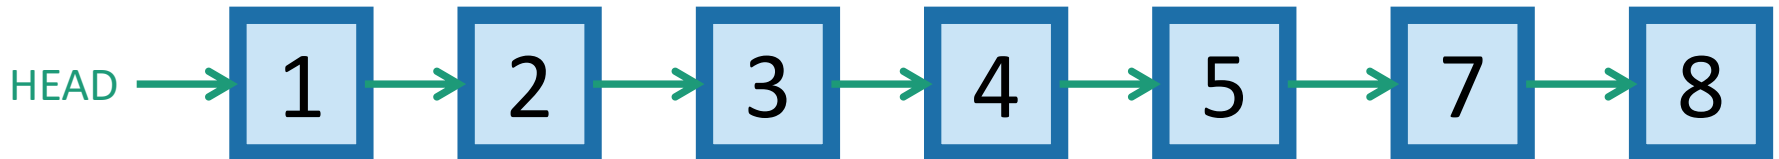
1. The punchline is **important**:
 - A data structure with $O(\log(n))$
INSERT/DELETE/SEARCH
2. The idea behind AVL Trees and **Red-Black Trees** is clever
 - It's good to be exposed to clever ideas.
 - Also it's just aesthetically pleasing.

Some data structures
for storing objects like **5** (aka, **nodes** with **keys**)

- (Sorted) arrays:



- (Sorted) linked lists:



- Some basic operations:
 - **INSERT, DELETE, SEARCH**

Sorted Arrays

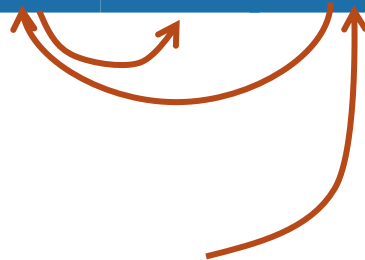
1	2	3	4	5	7	8
---	---	---	---	---	---	---

- $O(n)$ INSERT/DELETE:

1	2	3	4	4.5	7	8
---	---	---	---	-----	---	---

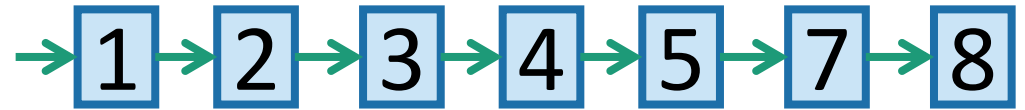
- $O(\log(n))$ SEARCH:

1	2	3	4	5	7	8
---	---	---	---	---	---	---



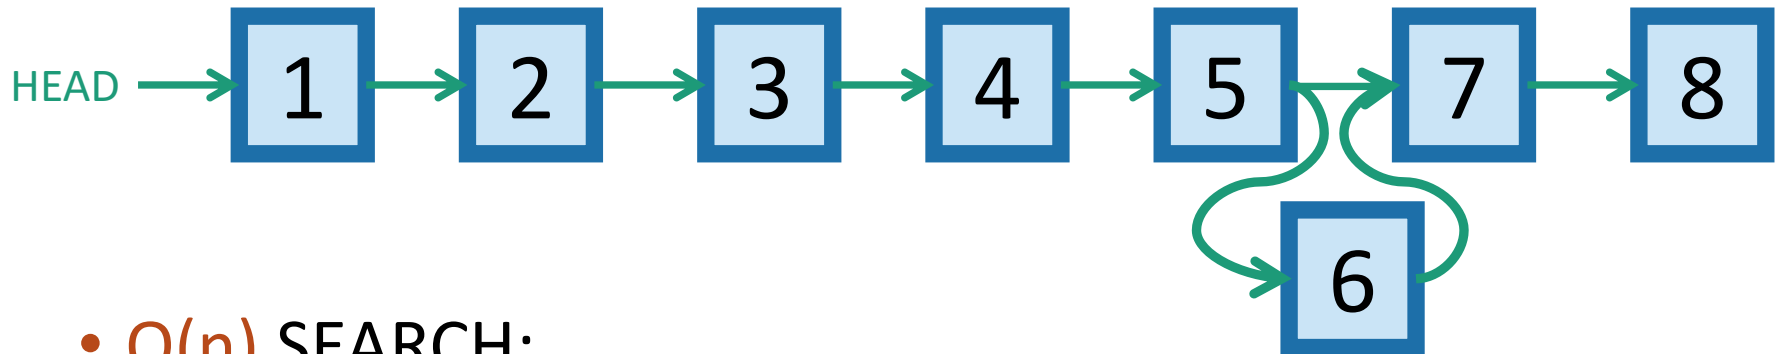
eg, Binary search to see if 3 is in A.

Sorted linked lists

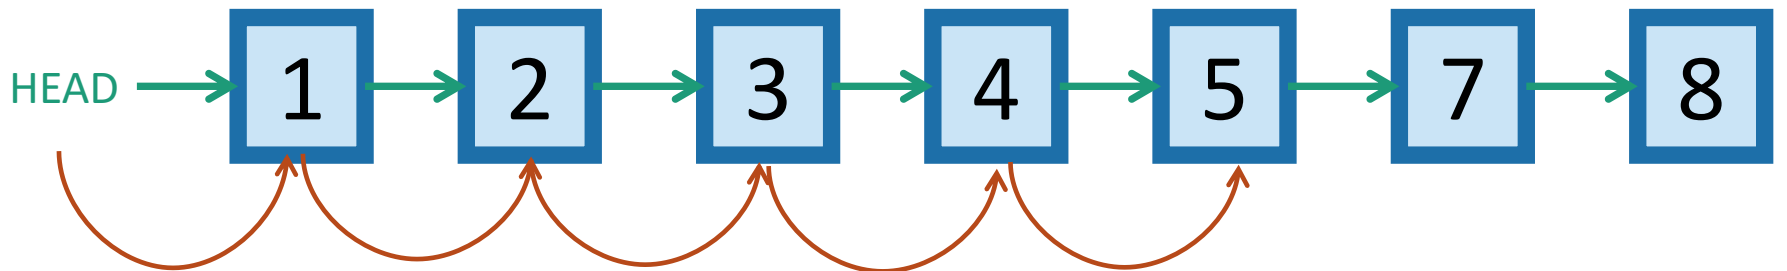


- $O(1)$ INSERT/DELETE:

- (assuming we have a pointer to the location of the insert/delete)









- $O(n)$ SEARCH:



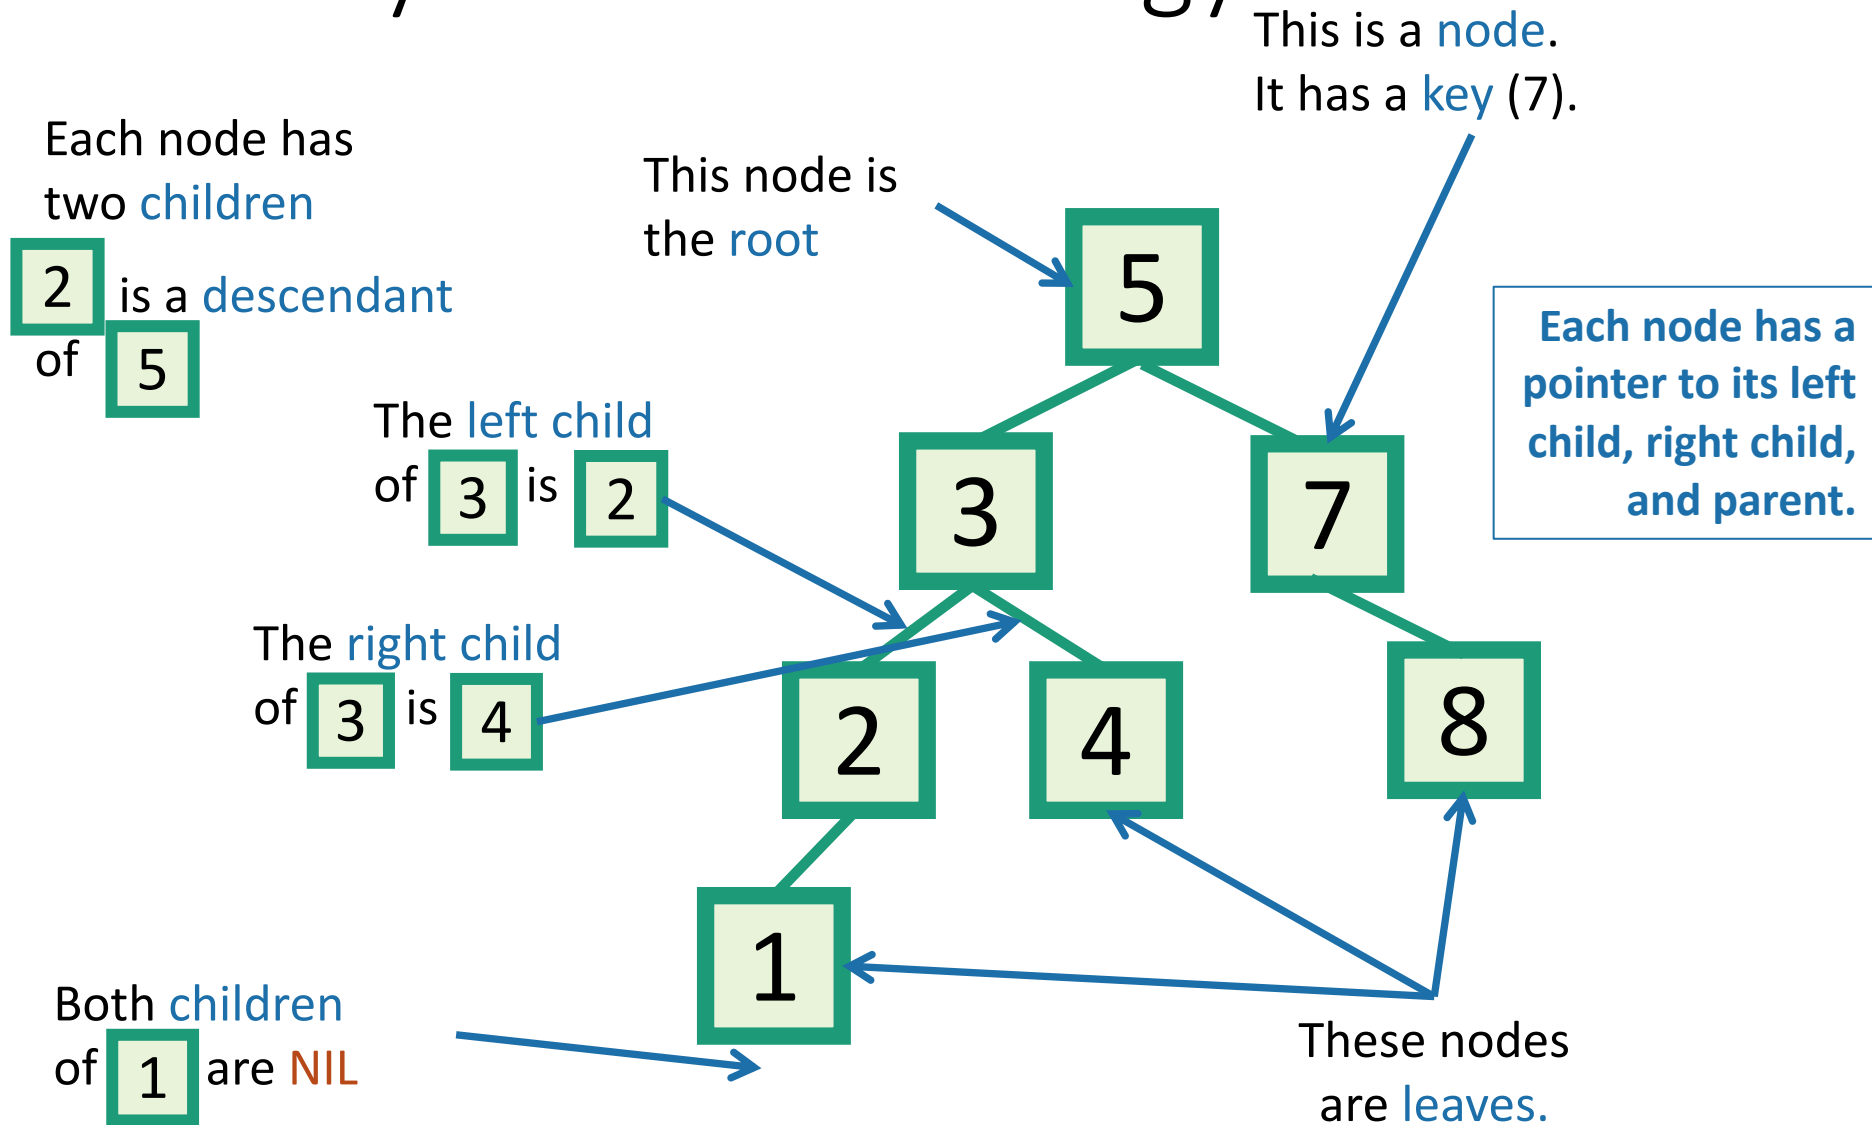
Motivation for Binary Search Trees

TODAY!

	Sorted Arrays	Linked Lists	Binary Search Trees*
Search	$O(\log(n))$ 	$O(n)$ 	$O(\log(n))$ 
Insert/Delete	$O(n)$ 	$O(1)$ 	$O(\log(n))$ 

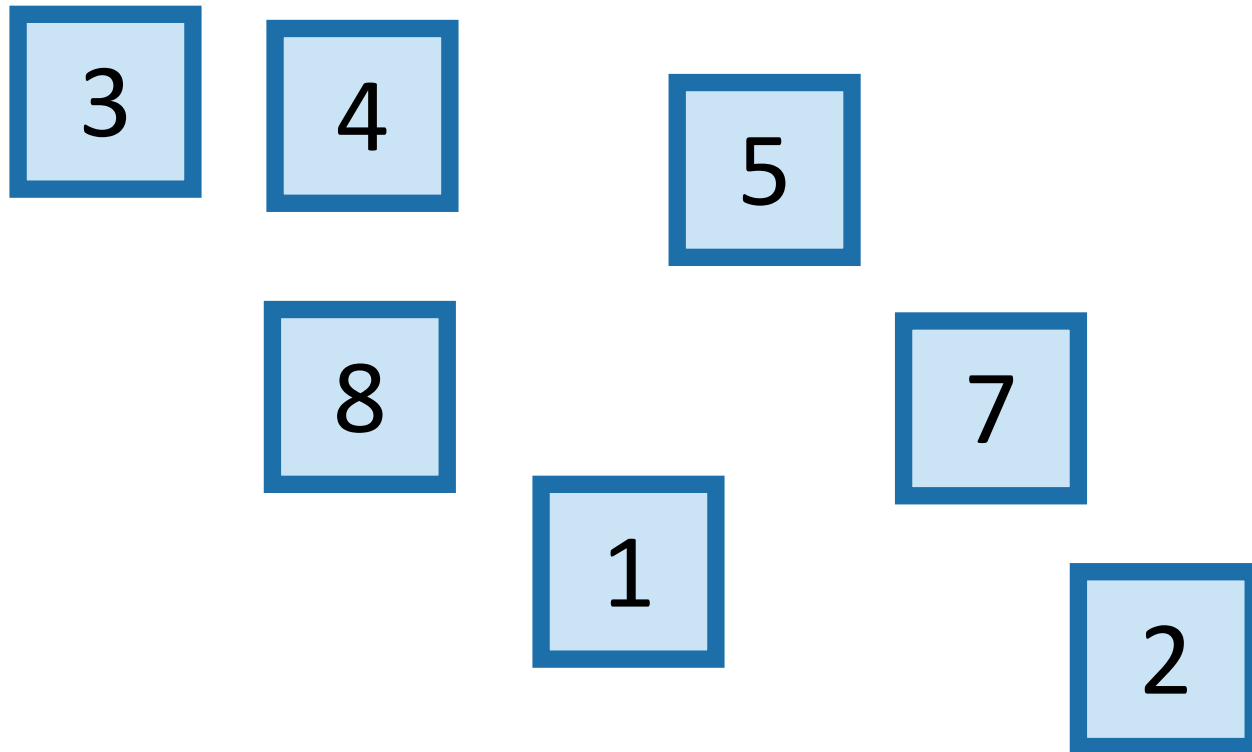
For today all keys are distinct.

Binary tree terminology



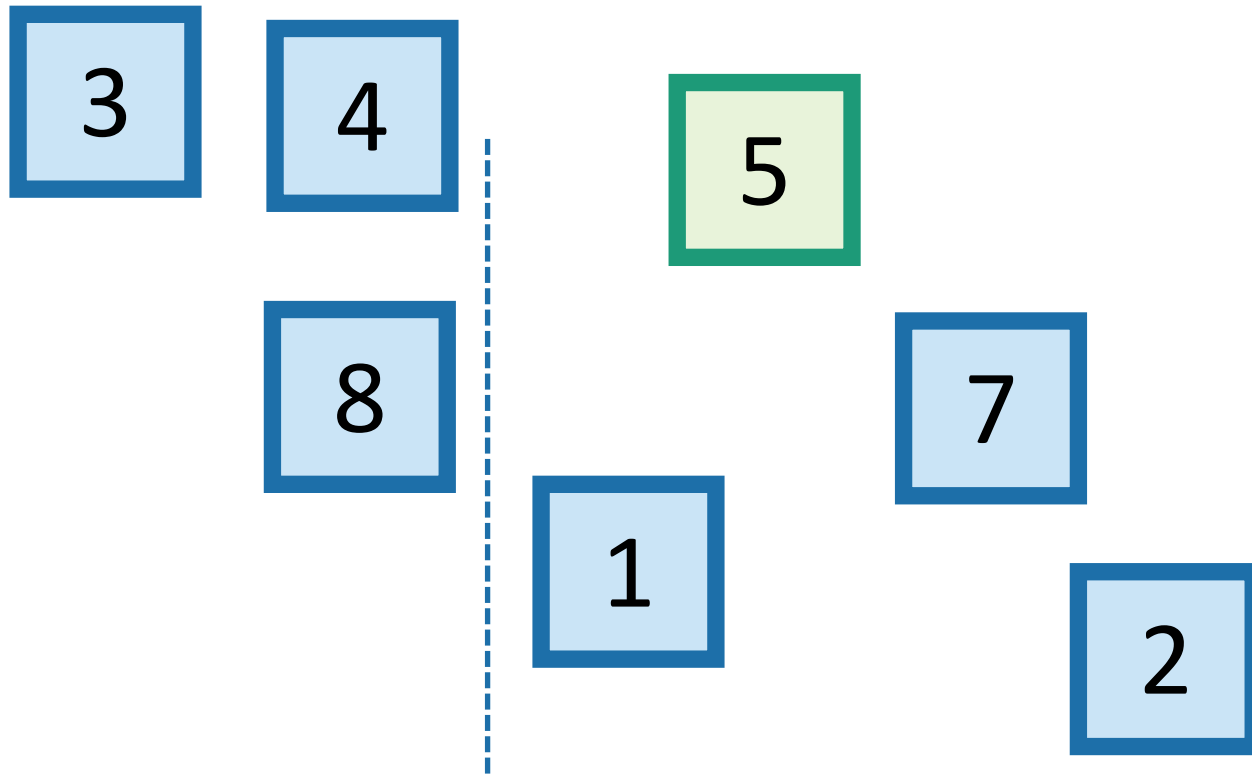
Binary Search Trees

- It's a **binary tree** so that:
 - Every LEFT descendant of a node has key less than that node.
 - Every RIGHT descendant of a node has key larger than that node.
- Example of building a binary search tree:



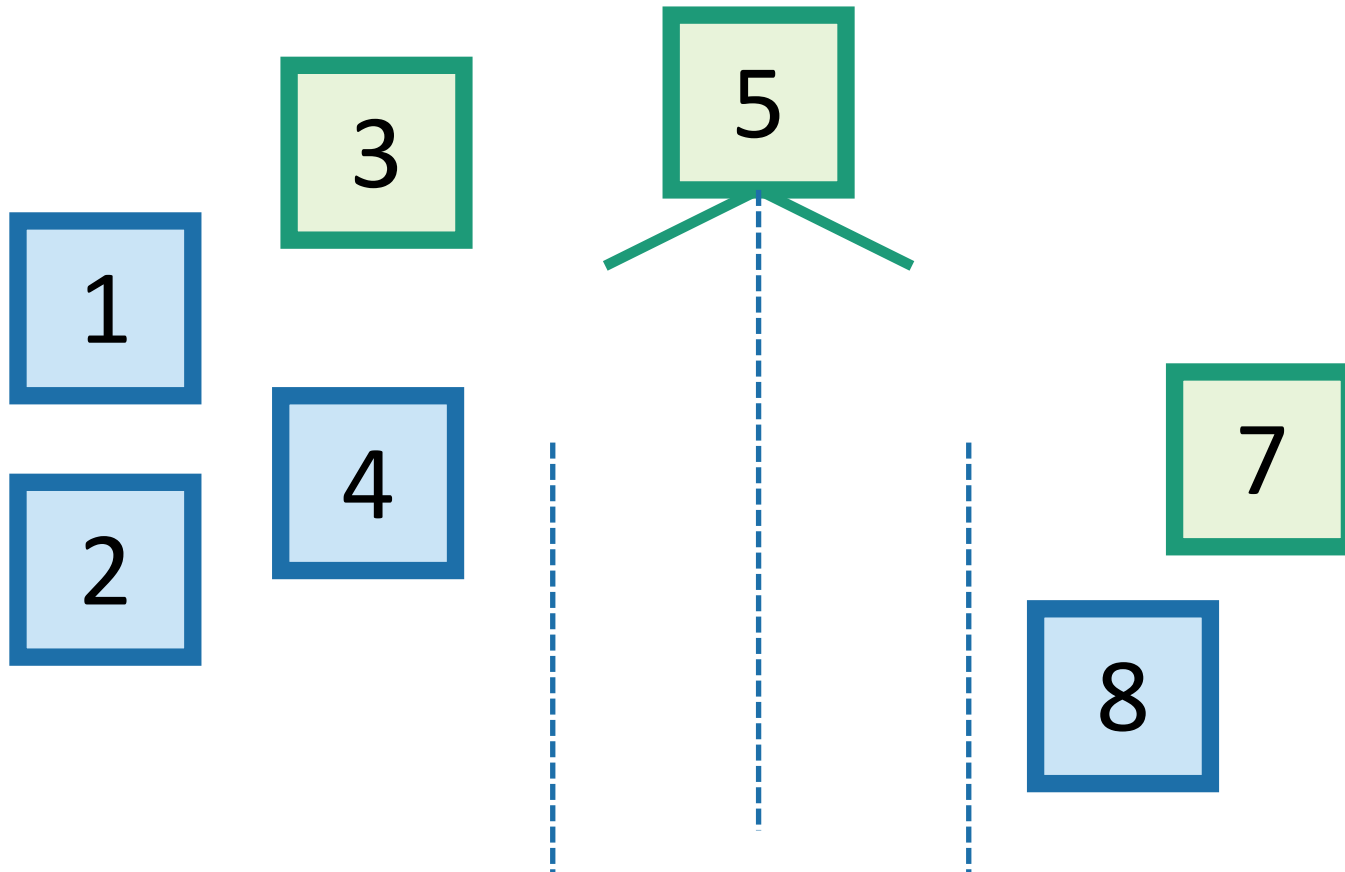
Binary Search Trees

- It's a **binary tree** so that:
 - Every LEFT descendant of a node has key less than that node.
 - Every RIGHT descendant of a node has key larger than that node.
- Example of building a binary search tree:




Binary Search Trees

- It's a **binary tree** so that:
 - Every LEFT descendant of a node has key less than that node.
 - Every RIGHT descendant of a node has key larger than that node.
- Example of building a binary search tree:



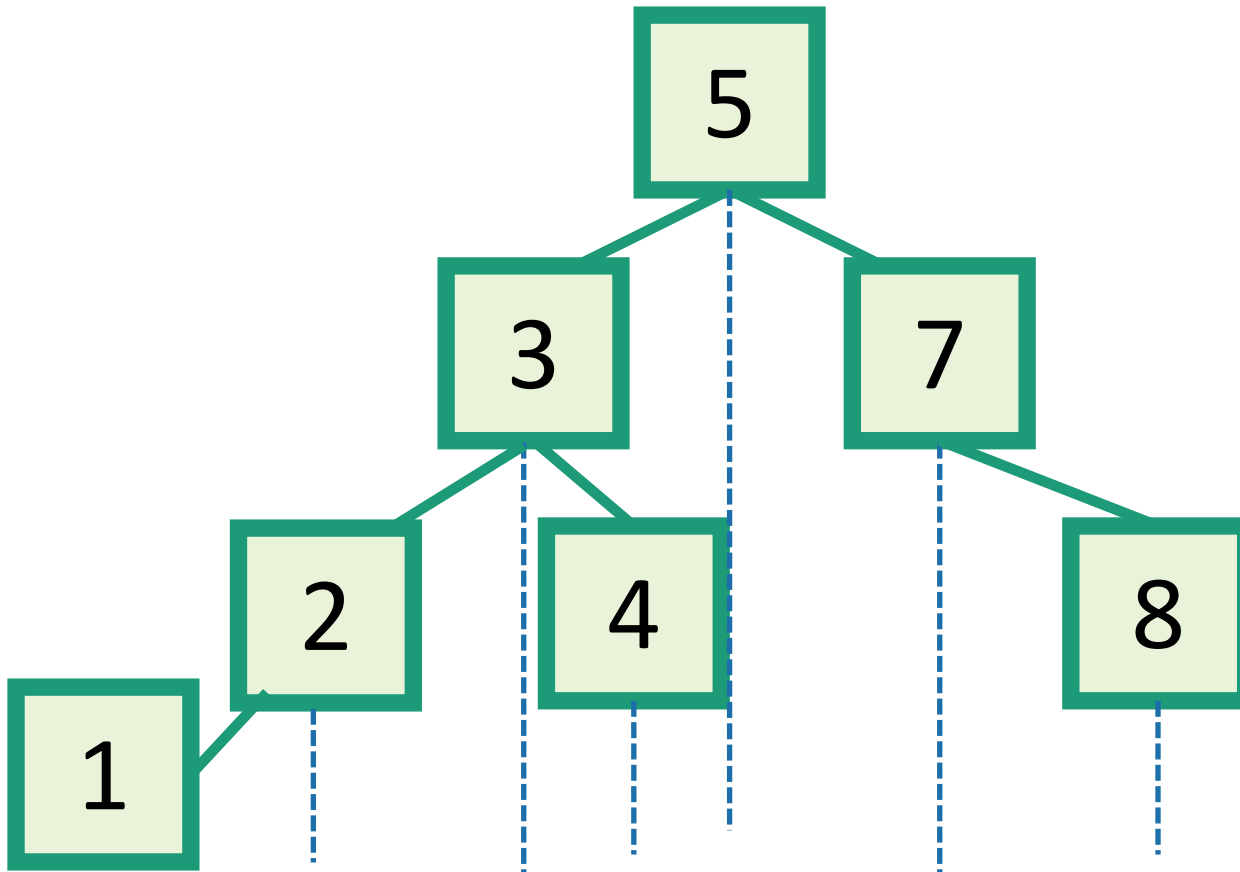
Binary Search Trees

- It's a **binary tree** so that:
 - Every LEFT descendant of a node has key less than that node.
 - Every RIGHT descendant of a node has key larger than that node.
- Example of building a binary search tree:



Binary Search Trees

- It's a **binary tree** so that:
 - Every LEFT descendant of a node has key less than that node.
 - Every RIGHT descendant of a node has key larger than that node.
- Example of building a binary search tree:



Q: Is this the only binary search tree I could possibly build with these values?

A: **No.** I made choices about which nodes to choose when. Any choices would have been fine.

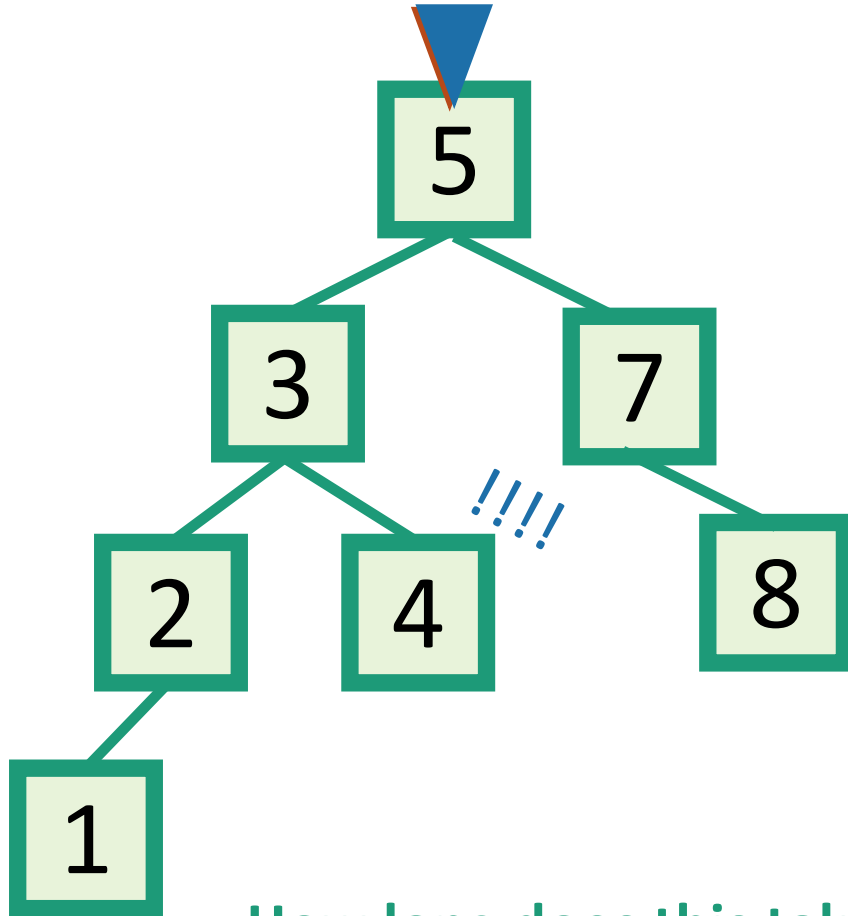
Remember the goal

Fast **SEARCH/INSERT/DELETE**

Can we do these?

SEARCH in a Binary Search Tree

definition by example



EXAMPLE: Search for 4.

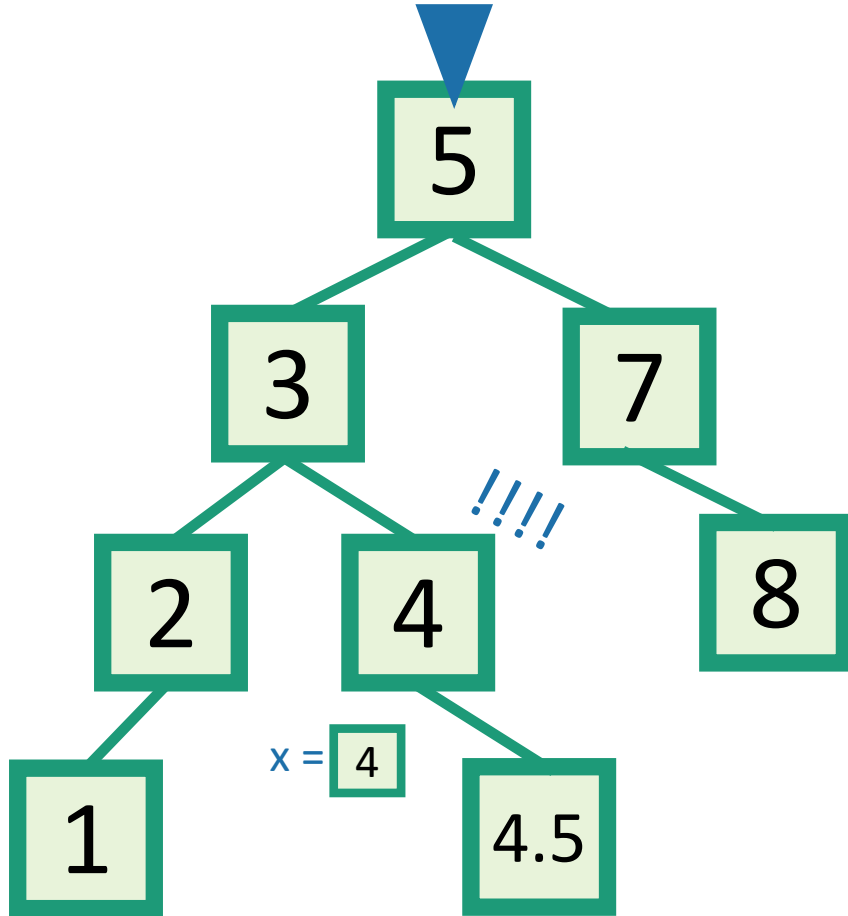
EXAMPLE: Search for 4.5

- It turns out it will be convenient to **return 4** in this case
- (that is, **return** the last node before we went off the tree)

How long does this take?

$O(\text{length of longest path}) = O(\text{height})$

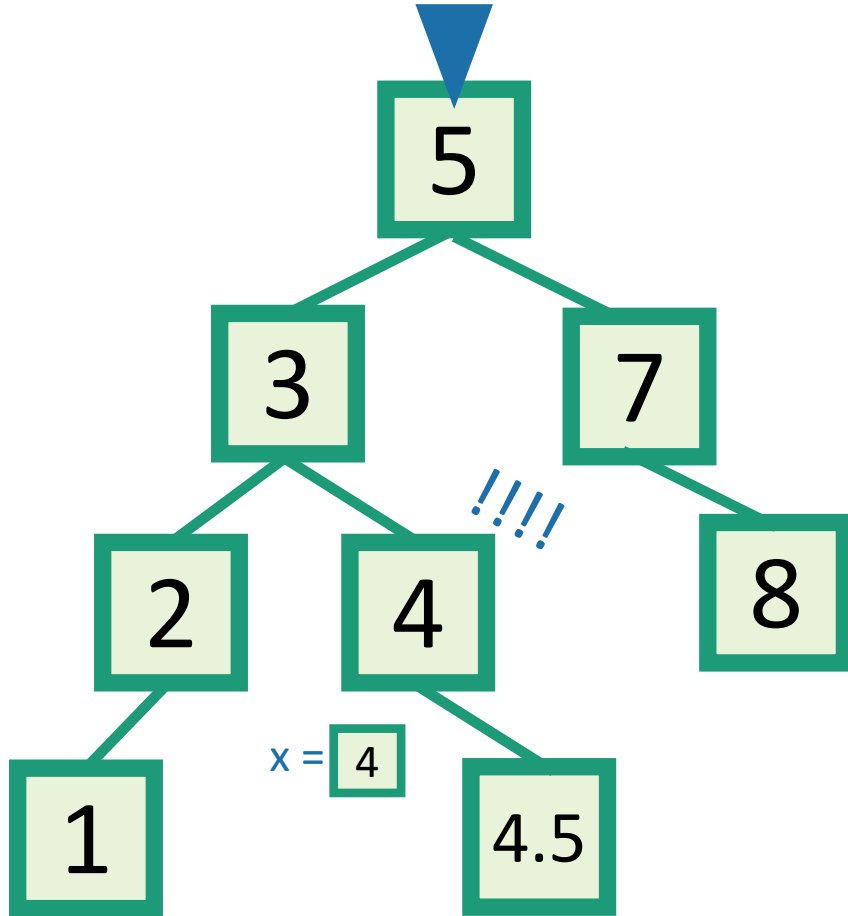
INSERT in a Binary Search Tree



EXAMPLE: Insert 4.5

- **INSERT**(key):
 - $x = \text{SEARCH}(\text{key})$
 - **Insert** a new node with desired key at x ...

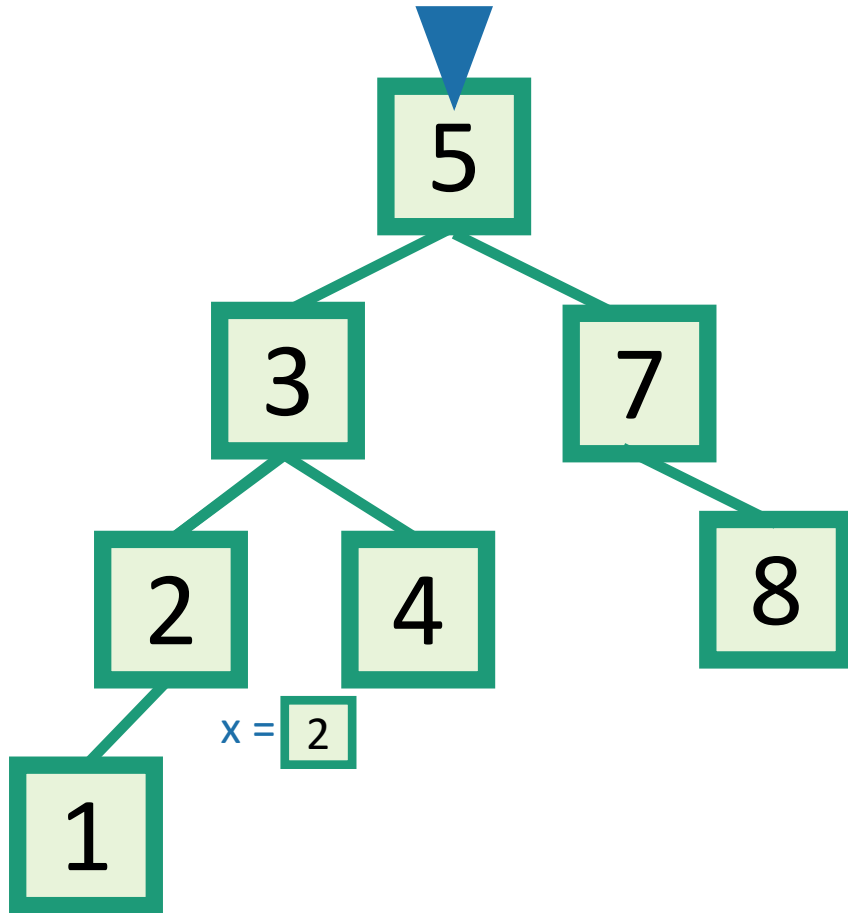
INSERT in a Binary Search Tree



EXAMPLE: Insert 4.5

- **INSERT**(key):
 - $x = \text{SEARCH}(\text{key})$
 - **if** $\text{key} > x.\text{key}$:
 - Make a new node with the correct key, and put it as the right child of x .
 - **if** $\text{key} < x.\text{key}$:
 - Make a new node with the correct key, and put it as the left child of x .
 - **if** $x.\text{key} == \text{key}$:
 - **return**

DELETE in a Binary Search Tree



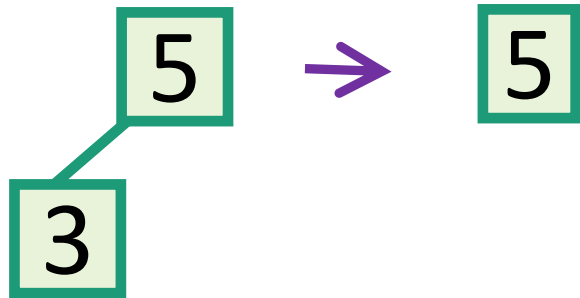
EXAMPLE: Delete 2

- **DELETE**(key):
 - $x = \text{SEARCH}(\text{key})$
 - **if** $x.\text{key} == \text{key}$:
 -delete x

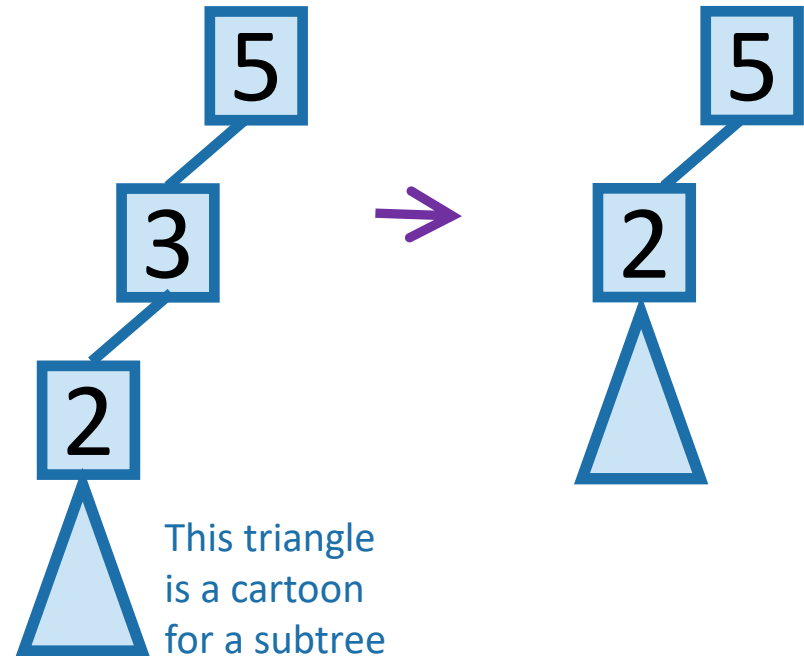
DELETE in a Binary Search Tree

several cases (by example)

say we want to delete 3



Case 1: if 3 is a leaf,
just delete it.




Case 2: if 3 has just one child,
move that up.

DELETE in a Binary Search Tree

ctd.

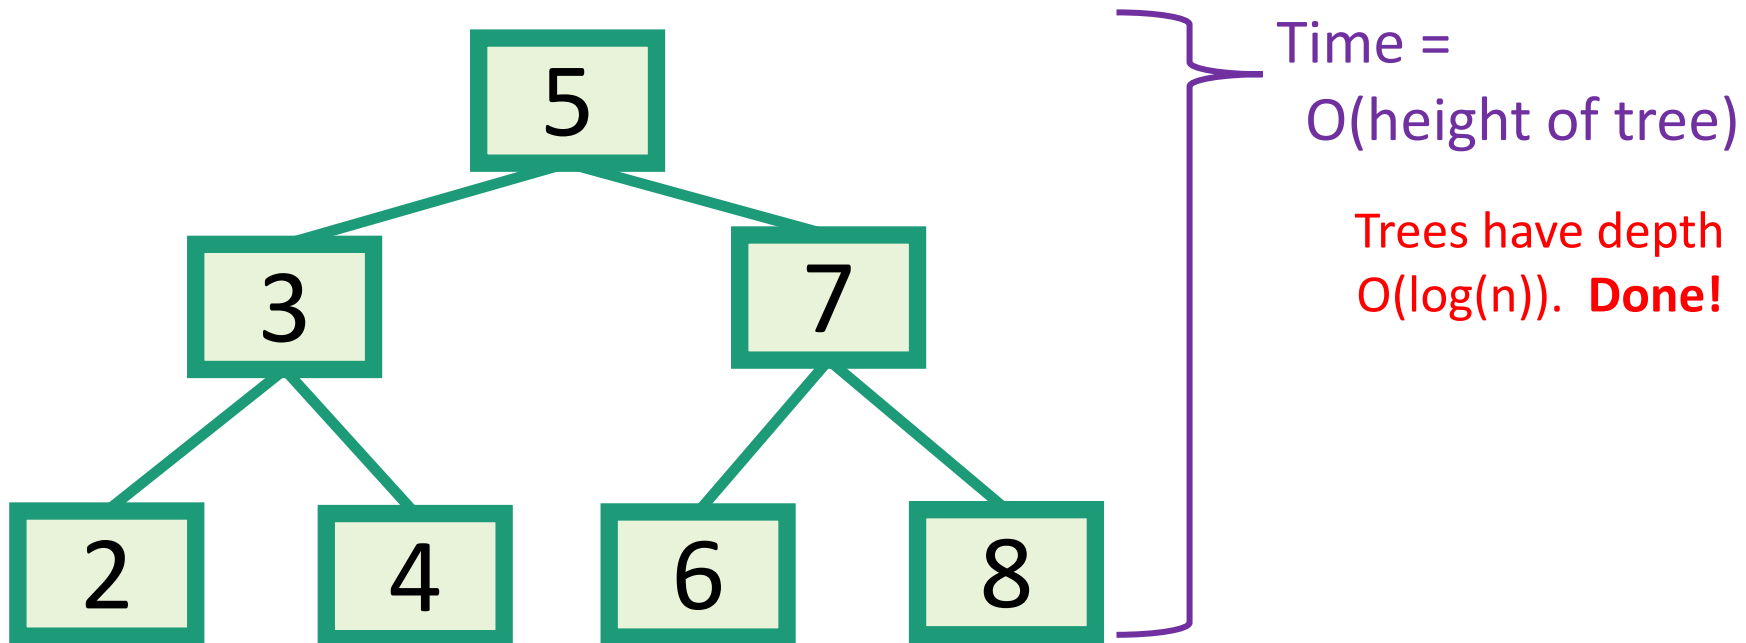
Case 3: if 3 has two children,
replace 3 with its **immediate successor**.
(aka, next biggest thing after 3)



- Does this maintain the BST property?
 - Yes.
- How do we find the immediate successor?
 - SEARCH for 3 in the subtree under 3.right
- How do we remove it when we find it?
 - If [3.1] has 0 or 1 children, do one of the previous cases.
- What if [3.1] has two children?
 - It doesn't.

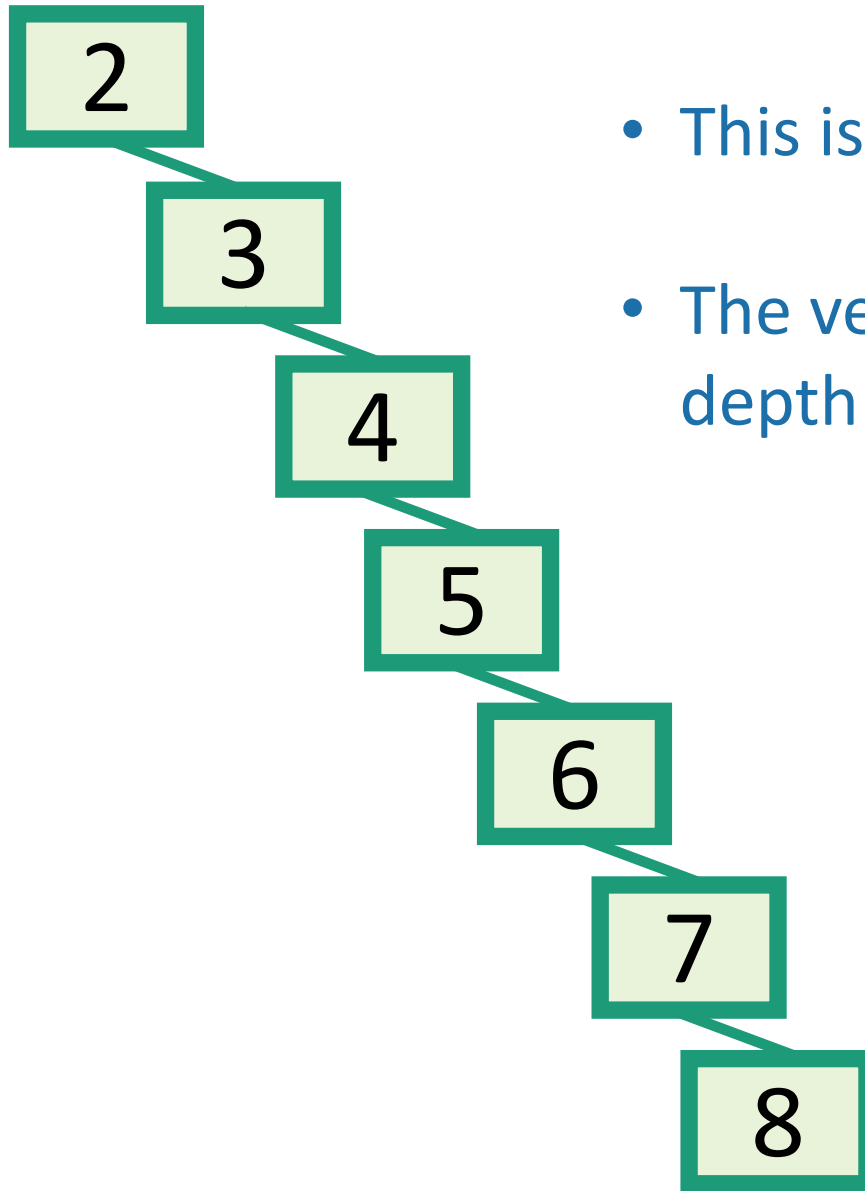
How long do these operations take?

- **SEARCH** is the big one.
 - Everything else just calls **SEARCH** and then does some small $O(1)$ -time operation.



How long does search take?

Wait...



- This is a valid binary search tree.
- The version with n nodes has depth n , **not** $O(\log(n))$.

Could such a tree show up?
In what order would I have to
insert the nodes?


Inserting in the order
2,3,4,5,6,7,8 would do it.



So this **could** happen.

What to do?

- Goal: Fast **SEARCH/INSERT/DELETE**
- All these things take time $O(\text{height})$
- And the height might be big!!! ☹️
- Idea 0:
 - Keep track of how deep the tree is getting.
 - If it gets too tall, re-do everything from scratch.
- Turns out that's not a great idea. Instead we turn to...

Self-Balancing Binary Search Trees


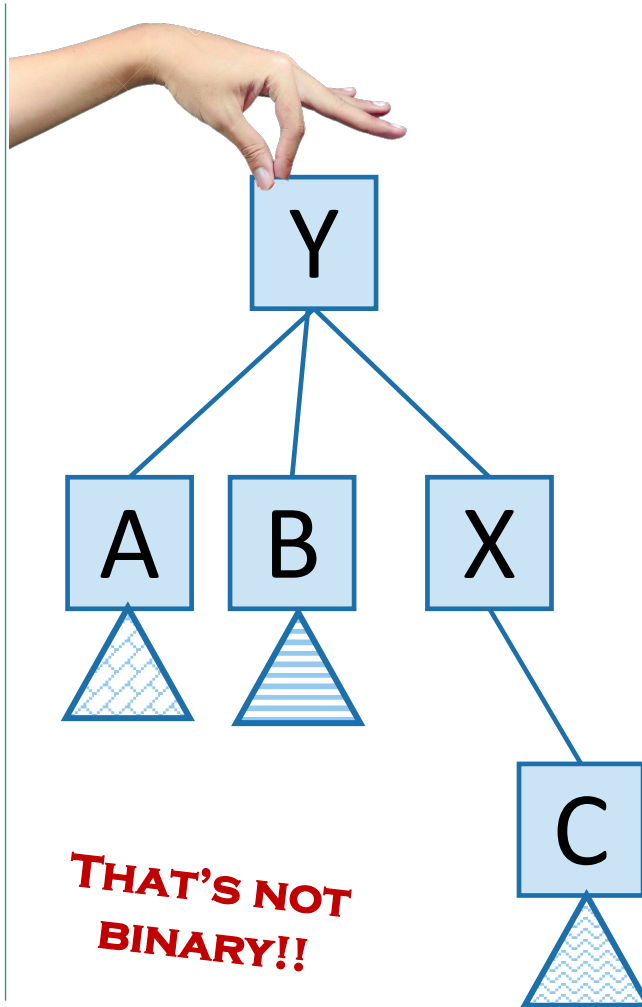





Idea 1: Rotations


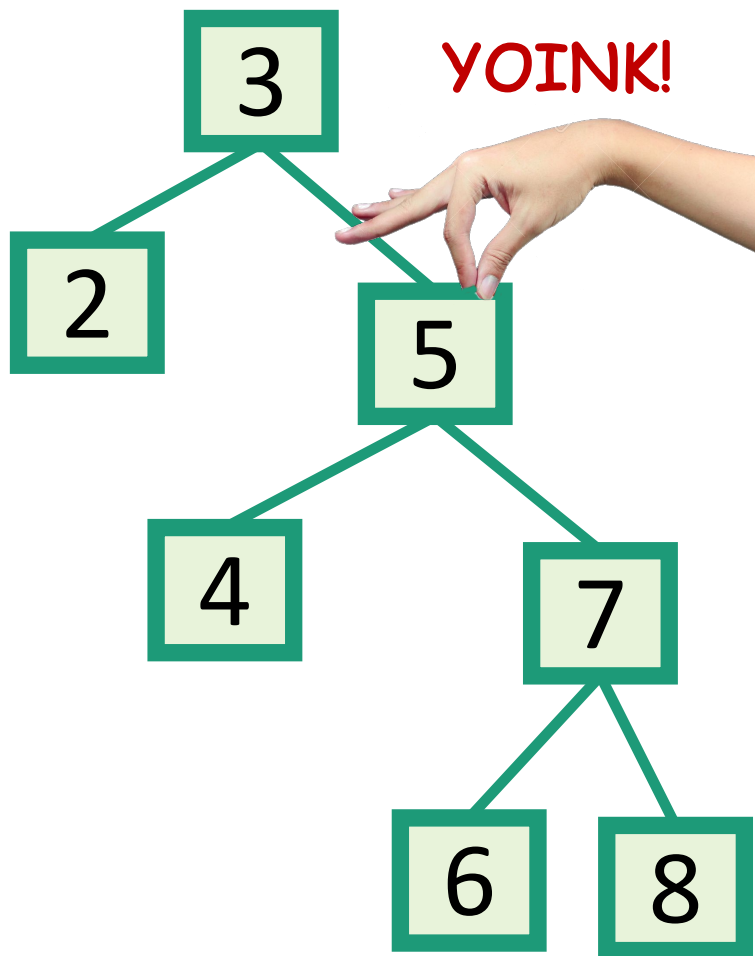
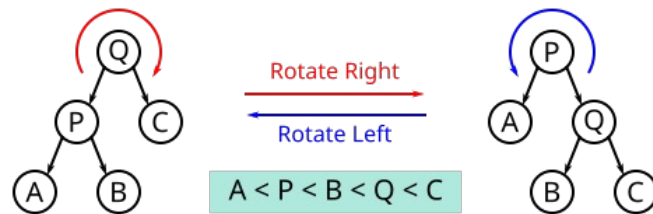
No matter what lives underneath A,B,C,
this takes time $O(1)$. (Why?)

- Maintain Binary Search Tree (BST) property, while moving stuff around.



Note: A, B, C, X, Y are variable names, not the contents of the nodes.

This seems helpful




Does this work?

- Whenever something seems unbalanced, do rotations until it's okay again.

Idea 2: have some proxy for balance



- Maintaining **perfect balance** is too hard.
- Instead, come up with some **proxy for balance**:
 - If the tree satisfies **[SOME PROPERTY]**, then it's pretty balanced.
 - We can maintain **[SOME PROPERTY]** using rotations.


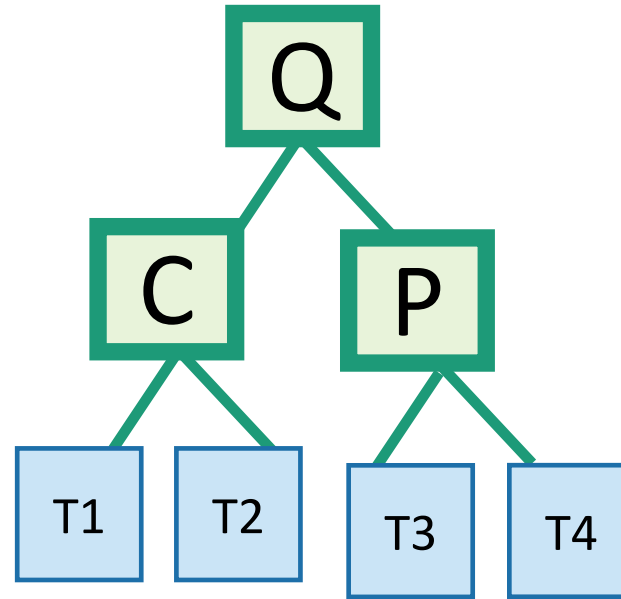
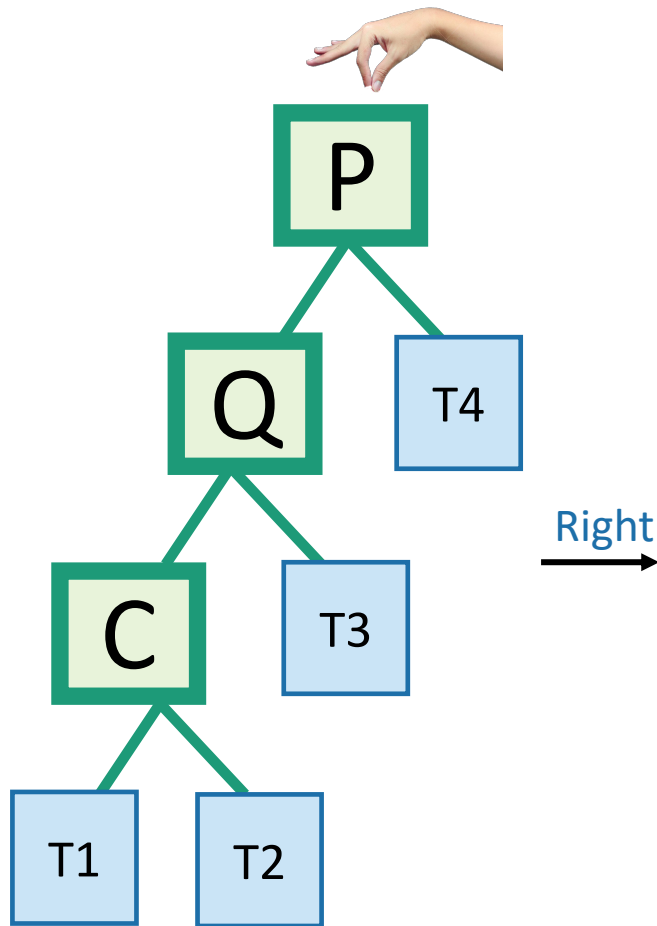


There are actually several ways to do this, but today we'll see...




AVL Trees

- Height-Balancing Property



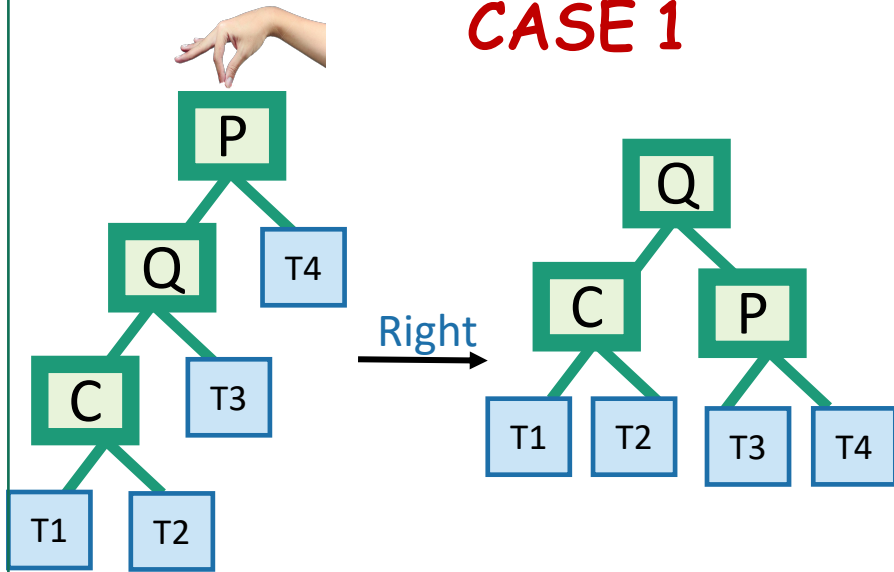


CASE 1

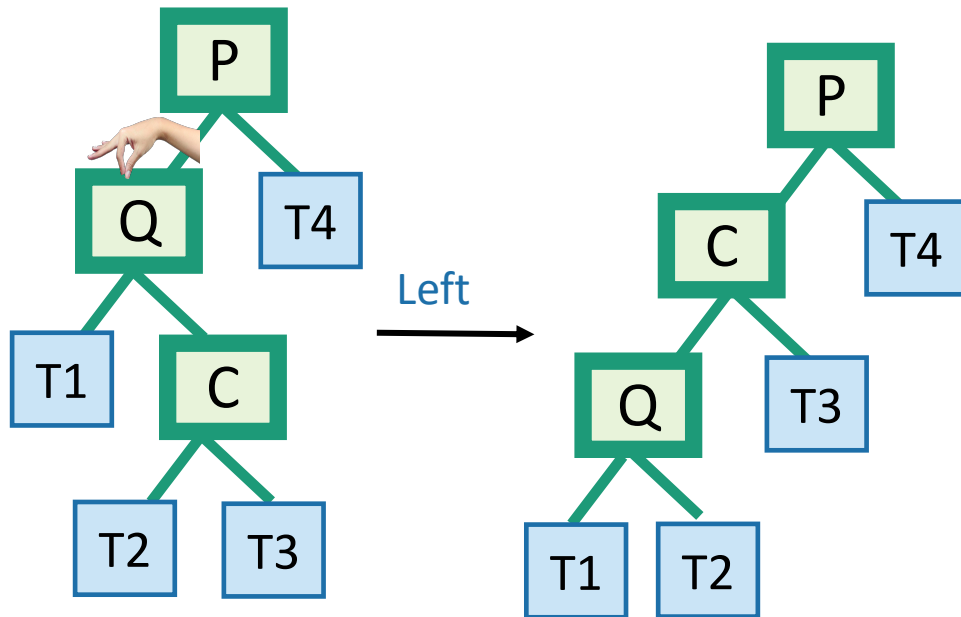
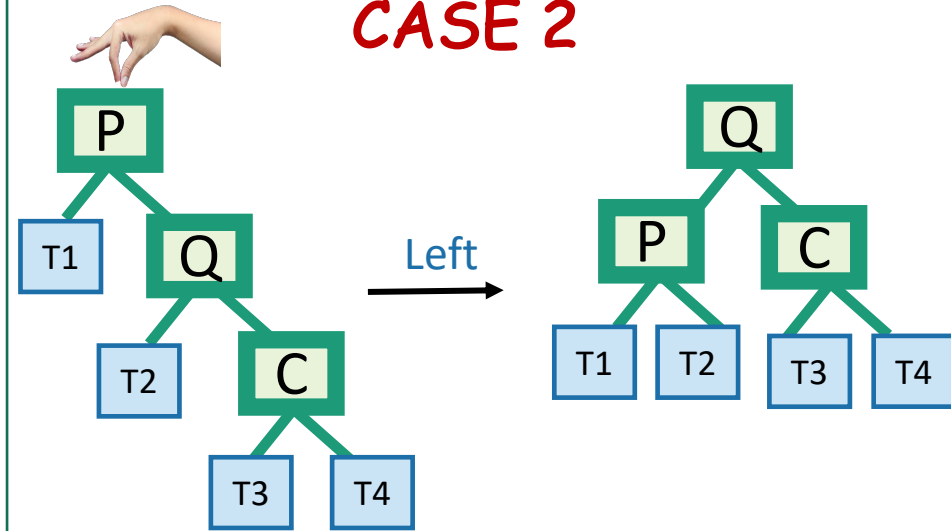


CASE 2

CASE 1

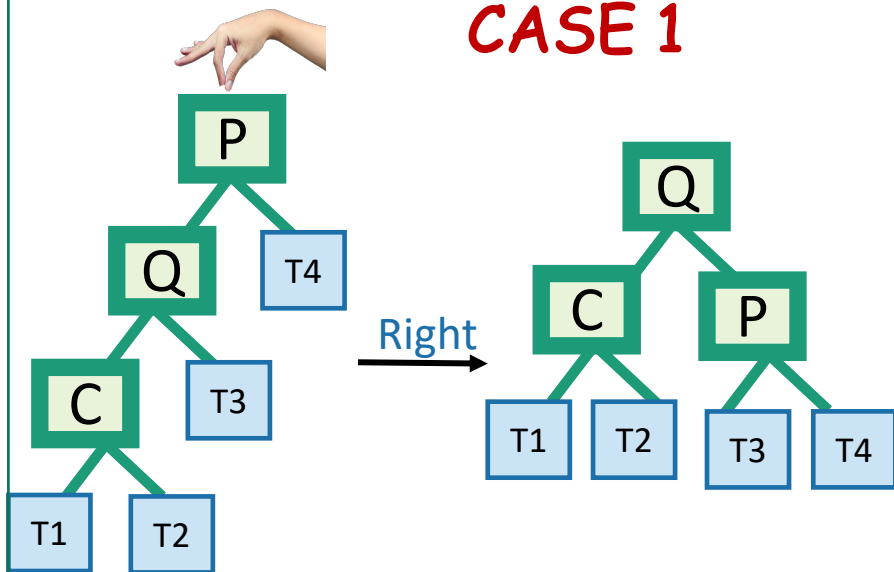


CASE 2

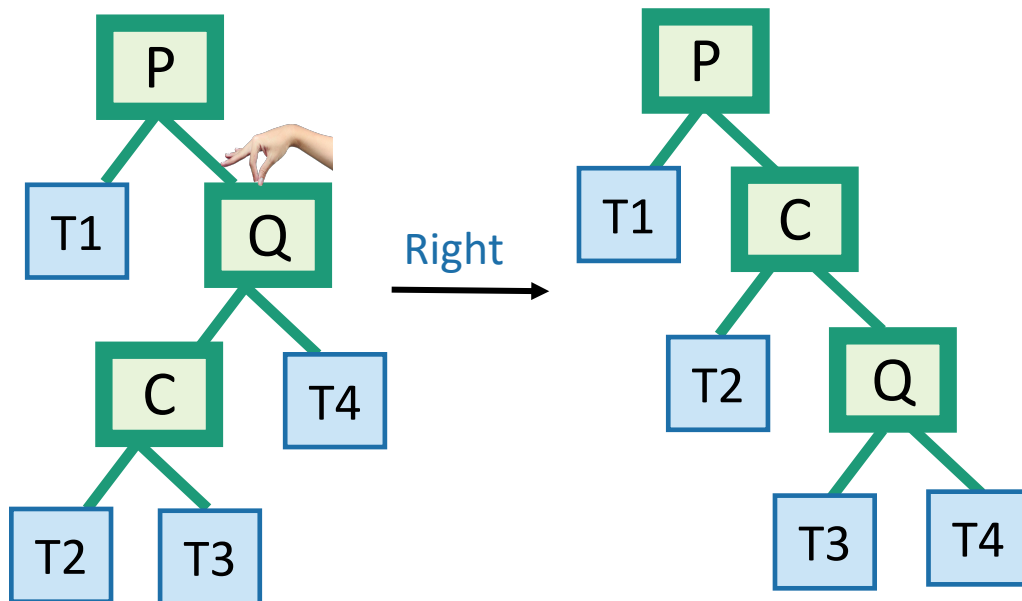
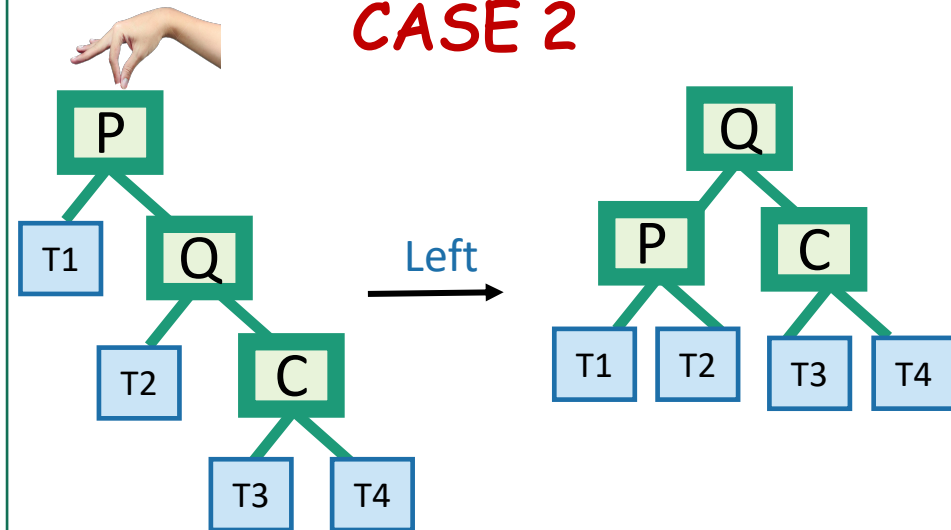


CASE 3

CASE 1

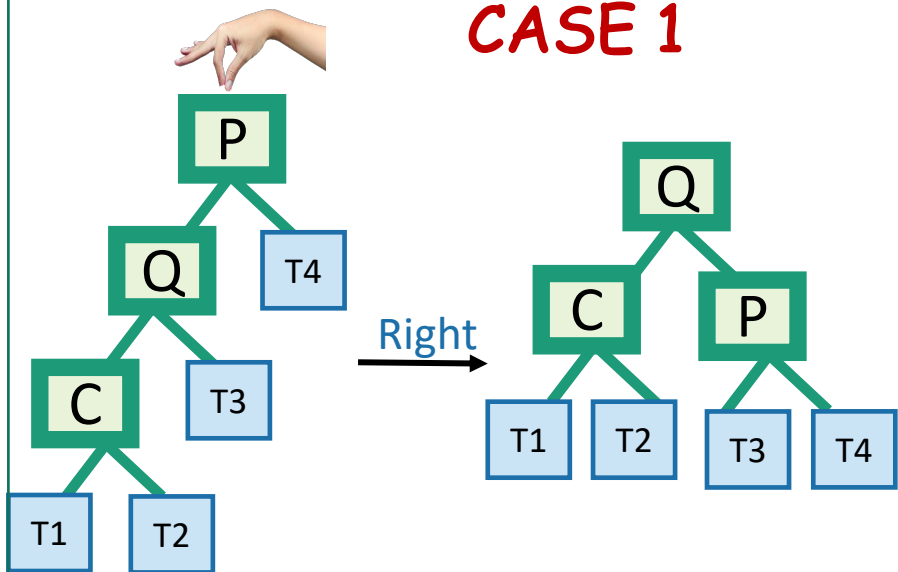


CASE 2




CASE 4

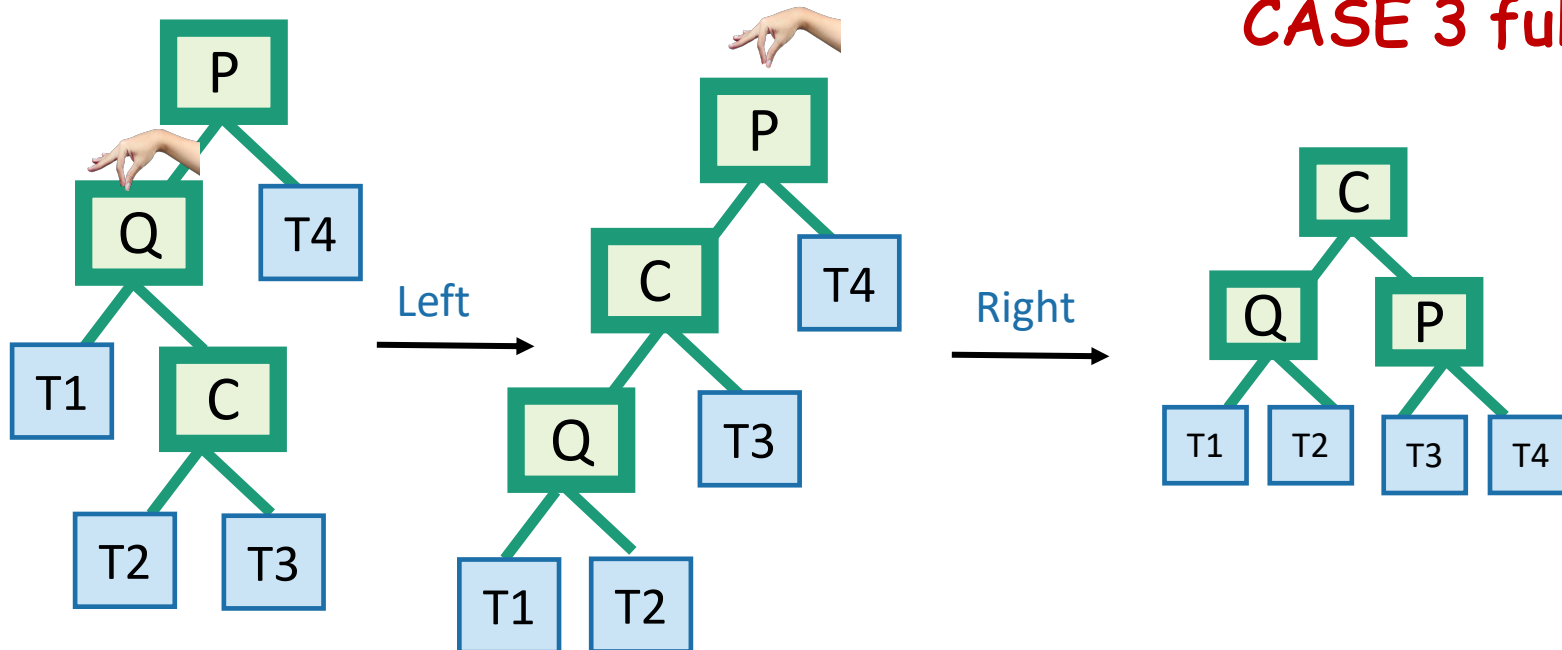
CASE 1



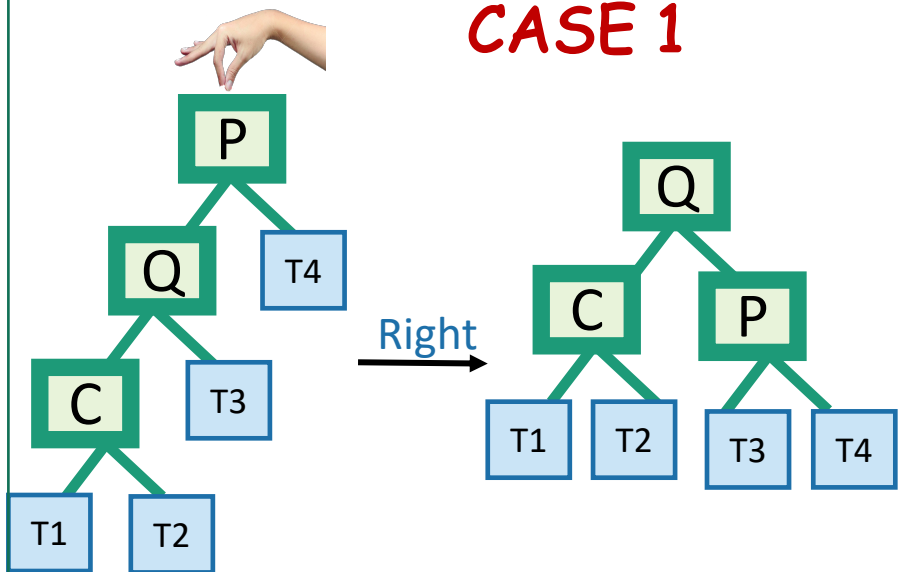
CASE 2



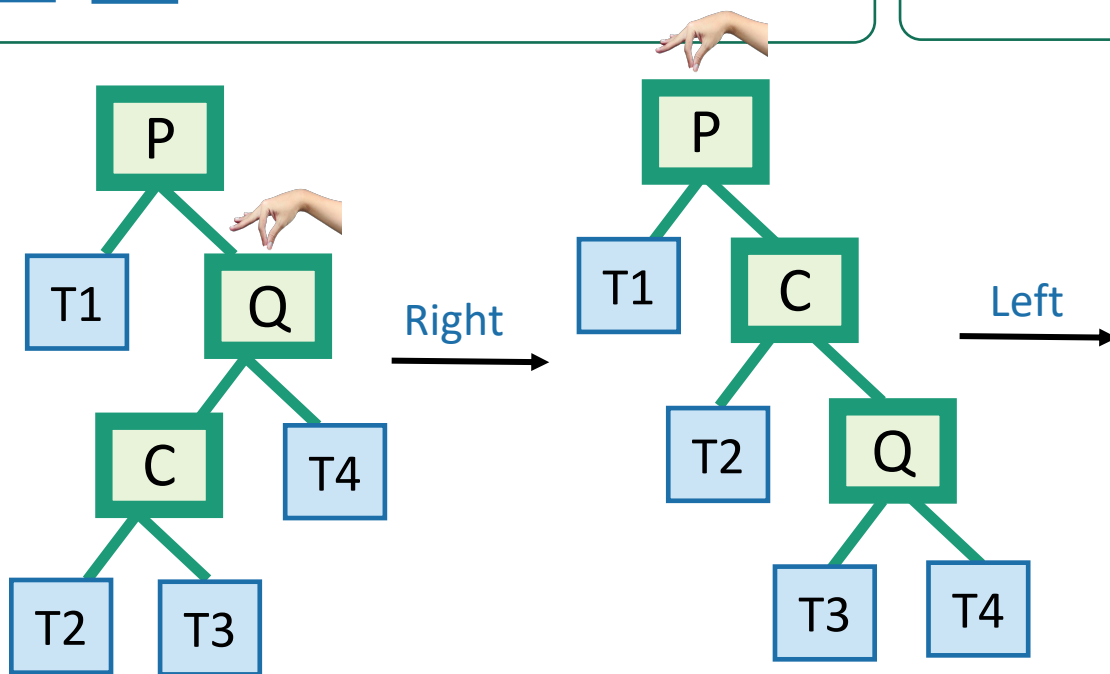

CASE 3 full



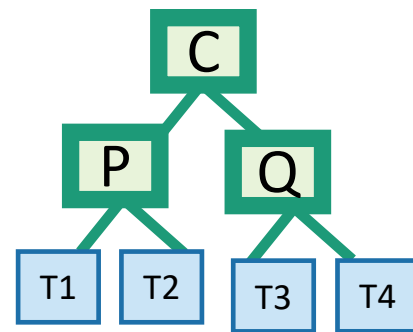
CASE 1



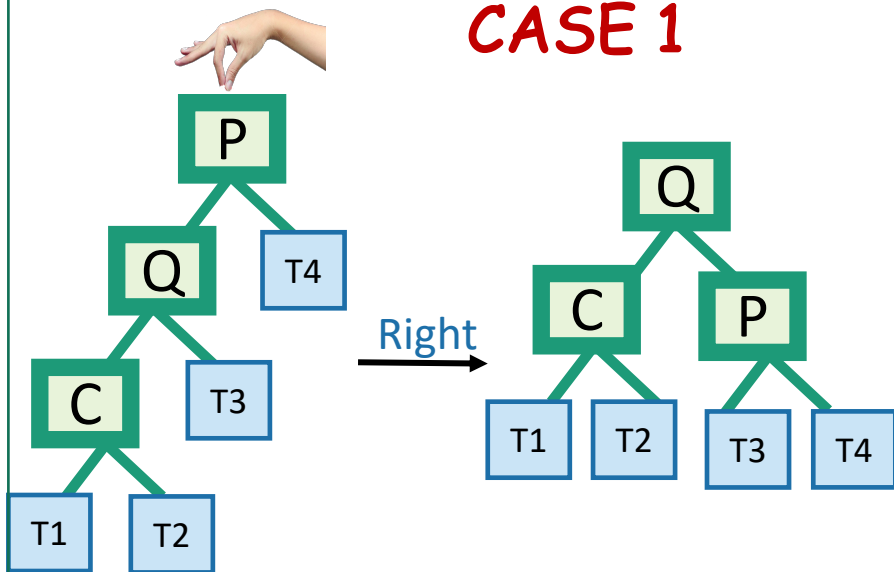
CASE 2



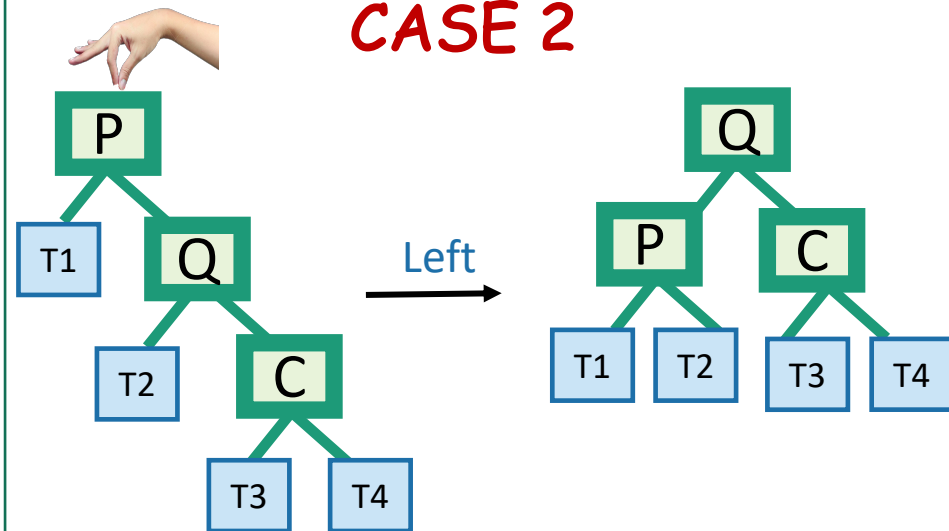
CASE 4 full



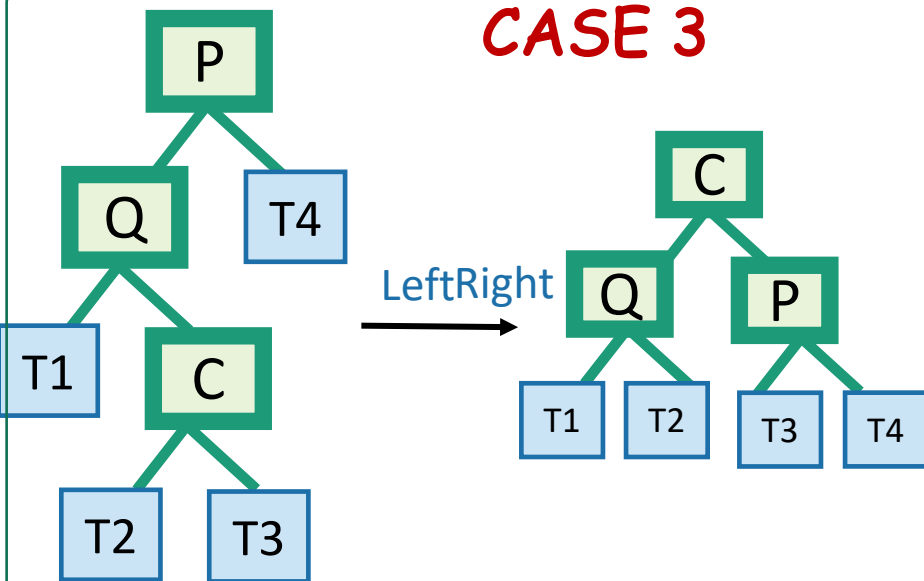
CASE 1



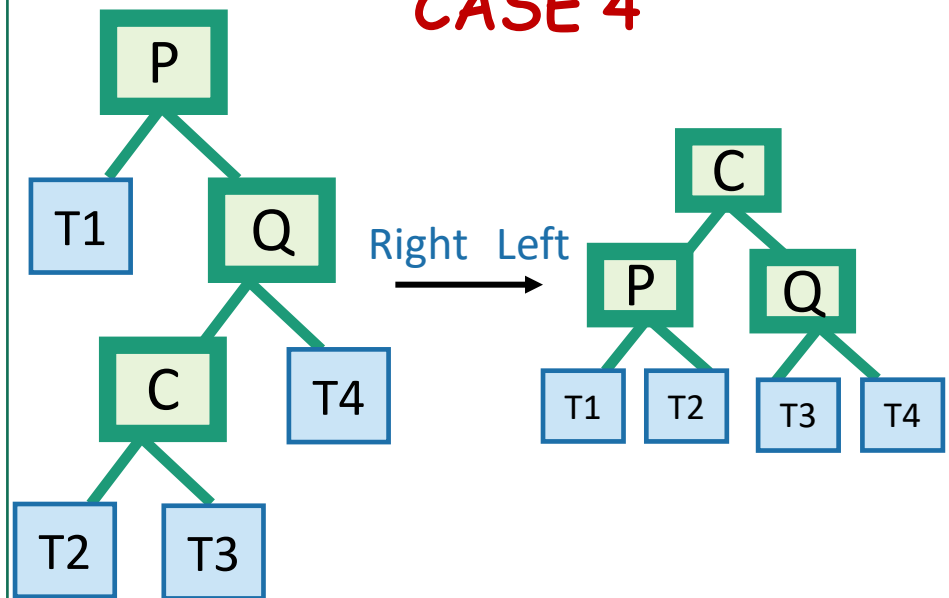
CASE 2



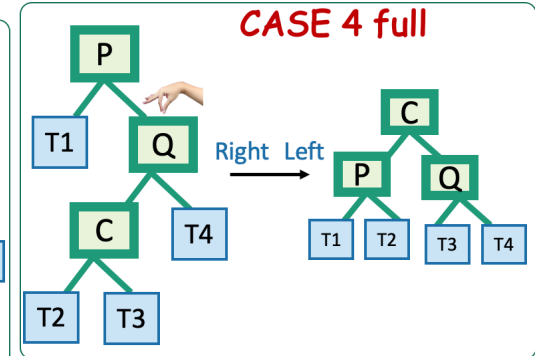
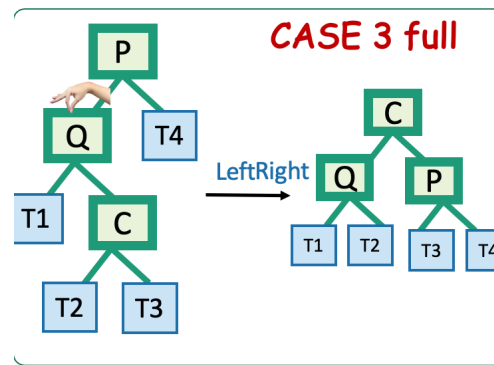


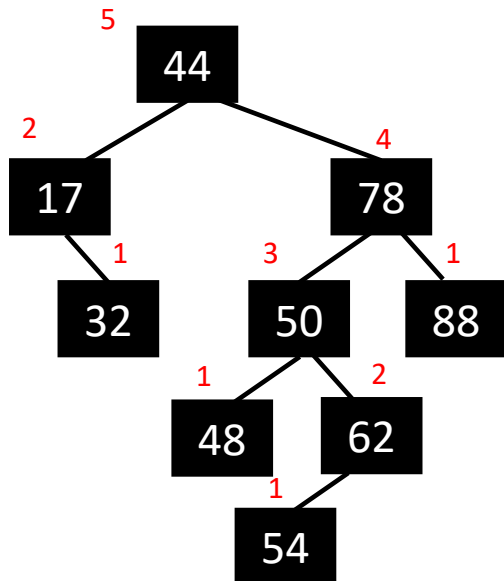
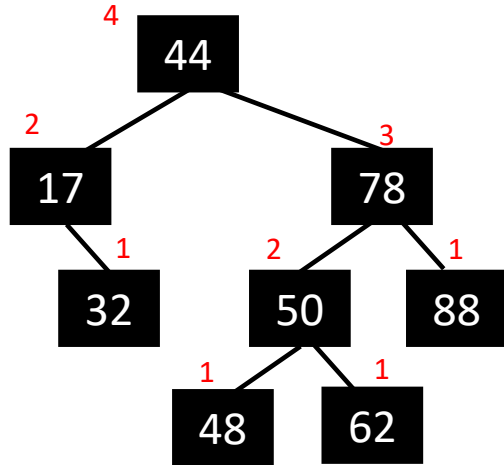
CASE 3





CASE 4




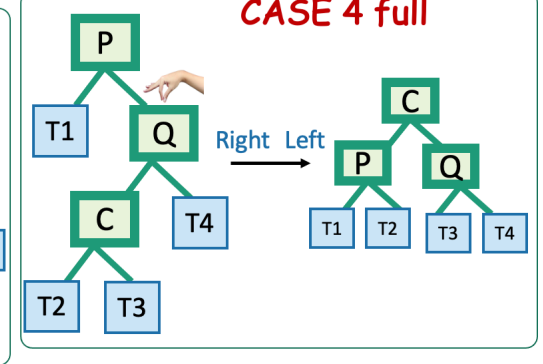




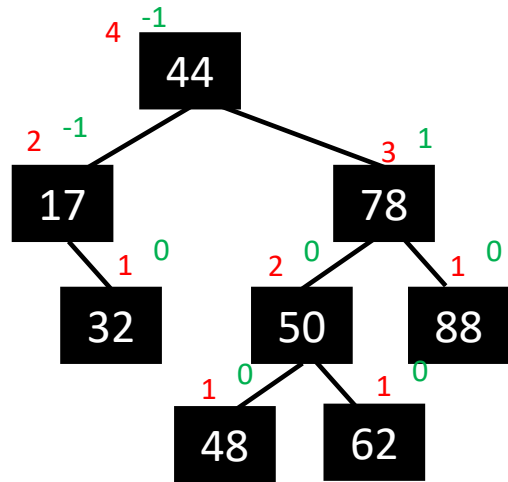
Inserting into an AVL Tree



Inserting into an AVL Tree


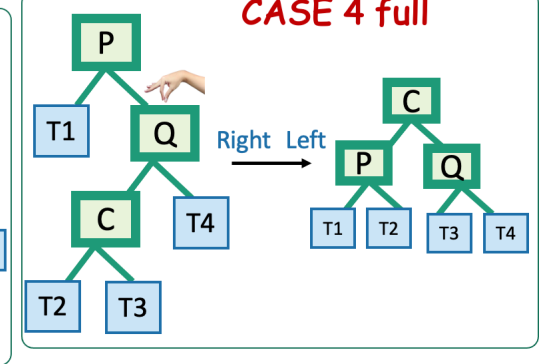
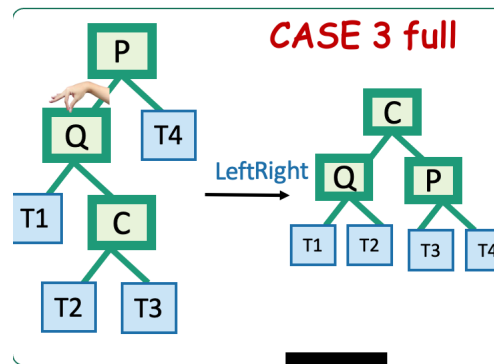





Inserting into an AVL Tree






Inserting into an AVL Tree

- **P** is the first unbalanced node you encountered going up from **p**, where **p** is the inserted node
- **Q** is the child of **P** with larger height (must be ancestor of **p**)
- **C** is the child of **Q** with larger height



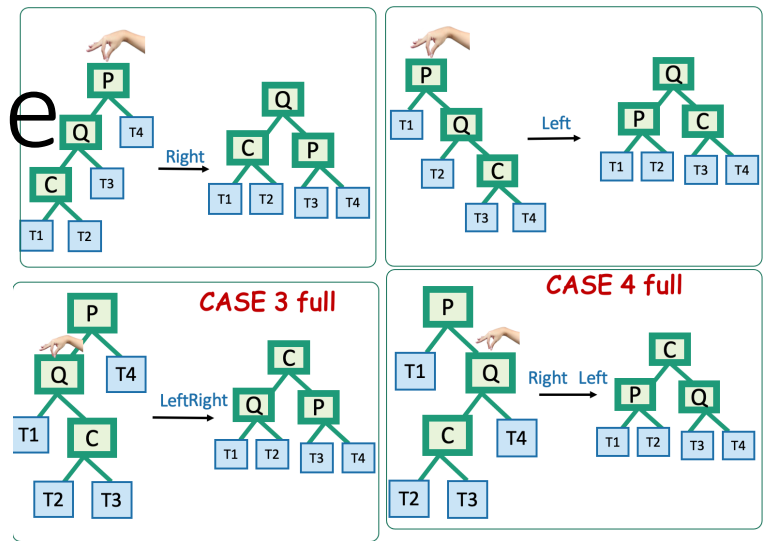


Deletion in an AVL Tree



Deletion in an AVL Tree

- **P** is the unbalanced first node you encountered going up from p, where p is the parent of removed node
- **Q** is the child of P with larger height
- **C** is the child of Q with larger height (if equal: try growing on the same side where P-Q are growing)



Red-Black Trees

- A Binary Search Tree that balances itself!
- No more time-consuming by-hand balancing!
- Be the envy of your friends and neighbors with the time-saving...

Red-Black tree!

Maintain balance by stipulating that **black nodes** are balanced, and that there aren't too many **red nodes**.


It's just good sense!



Red-Black Trees


these rules are the proxy for balance

- Every node is colored **red** or **black**.
- The root node is a **black node**.
- NIL children count as **black nodes**.
- Children of a **red node** are **black nodes**.
- For all nodes x :
 - all paths from x to NIL's have the same number of **black nodes** on them.






I'm not going to draw the NIL children in the future, but they are treated as black nodes.

Examples(?)




- Every node is colored **red** or **black**.
- The root node is a **black node**.
- NIL children count as **black nodes**.
- Children of a **red node** are **black nodes**.
- For all nodes x:
 - all paths from x to NIL's have the same number of **black nodes** on them.



Why??????

- This is pretty balanced.
 - The **black nodes** are balanced
 - The **red nodes** are “spread out” so they don’t mess things up too much.
- We can maintain this property as we insert/delete nodes, by using rotations.




This is the really clever idea!

This **Red-Black** structure is a proxy for balance.

This is “pretty balanced”

- To see why, intuitively, let's try to build a Red-Black Tree that's **unbalanced**.



Conjecture: the height of a
red-black tree is at most $2 \log(n)$

Okay, so it's balanced...




...but can we maintain it?

- Yes!



- See CLRS for more details.

- (You are not responsible for the details for this class – but you should understand the main ideas).





CASE 1



CASE 2



CASE 3





RB-INSERT-FIXUP(T, z)



```
1  while  $z.p.color == \text{RED}$ 
2      if  $z.p == z.p.p.left$ 
3           $y = z.p.p.right$ 
4          if  $y.color == \text{RED}$ 
5               $z.p.color = \text{BLACK}$  // case 1
6               $y.color = \text{BLACK}$  // case 1
7               $z.p.p.color = \text{RED}$  // case 1
8               $z = z.p.p$  // case 1
9          else if  $z == z.p.right$ 
10              $z = z.p$  // case 2
11             LEFT-ROTATE( $T, z$ ) // case 2
12              $z.p.color = \text{BLACK}$  // case 3
13              $z.p.p.color = \text{RED}$  // case 3
14             RIGHT-ROTATE( $T, z.p.p$ ) // case 3
15         else (same as then clause
16             with “right” and “left” exchanged)
```

16 $T.root.color = \text{BLACK}$

SE 1



SE 2



RB-INSERT-FIXUP(T, z)

```

1  while  $z.p.color == RED$ 
2      if  $z.p == z.p.p.left$ 
3           $y = z.p.p.right$ 
4          if  $y.color == RED$ 
5               $z.p.color = BLACK$  // case 1
6               $y.color = BLACK$  // case 1
7               $z.p.p.color = RED$  // case 1
8               $z = z.p.p$  // case 1
9          else if  $z == z.p.right$ 
10              $z = z.p$  // case 2
11             LEFT-ROTATE( $T, z$ ) // case 2
12              $z.p.color = BLACK$  // case 3
13              $z.p.p.color = RED$  // case 3
14             RIGHT-ROTATE( $T, z.p.p$ ) // case 3
15          else (same as then clause
              with "right" and "left" exchanged)
16   $T.root.color = BLACK$ 

```

Deleting from a Red-Black tree

Fun exercise!






That's a lot of cases

- You are **not responsible** for the nitty-gritty details of Red-Black Trees. (For this class)
 - Though implementing them is a great exercise!
- You should know:
 - What are the properties of an RB tree?
 - And (more important) why does that guarantee that they are balanced?

What was the point again?

- Red-Black Trees **always** have height at most $2\log(n+1)$.
- As with general **Binary Search Trees**, all operations are $O(\text{height})$
- So all operations are $O(\log(n))$.

Conclusion: The best of both worlds

	Sorted Arrays	Linked Lists	Balanced Binary Search Trees
Search	$O(\log(n))$ 	$O(n)$ 	$O(\log(n))$ 
Insert/Delete	$O(n)$ 	$O(1)$ 	$O(\log(n))$ 